

Lecture 21: Correlated Errors (Continued)

Ailin Zhang

2023-07-11

Agenda

- Definition of Correlated Errors
- Detection of Correlated Errors
 - Time plot of residuals
 - Runs test
 - Durbin-Watson test
 - Lag plots
 - Autocorrelation function and autocorrelation plot
- Remedies to Correct for Autocorrelated Errors (Today)
- Autocorrelation and Seasonality (Today)

Autocorrelation

- The lag-k autocorrelation of the residuals (e_1, \dots, e_n) is defined as

$$\hat{\rho}_k = \frac{\sum_{t=k+1}^n e_t e_{t-k}}{\sum_{t=1}^n e_t^2}$$

- Autocorrelation is slightly different from the correlation of (e_1, \dots, e_{n-k}) v.s. (e_{1+k}, \dots, e_n) :

$$\frac{\sum_{t=k+1}^n e_t e_{t-k}}{\sqrt{\sum_{t=1}^{n-k} e_t^2 \sum_{t=k+1}^n e_t^2}}$$

- The R command `acf()` (autocorrelation function) in R can calculate lag-k autocorrelation.


```
acf(lm1$res, lag.max = 5, plot=FALSE)
```

```
##
## Autocorrelations of series 'lm1$res', by lag
##
##      0      1      2      3      4      5
## 1.000 0.751 0.521 0.297 -0.007 -0.220
```

Autocorrelation Function and the Plot

In time-series analysis, one often plot the lag-k autocorrelations against k to examine the autocorrelation structure of a variable. The `acf()` command can produce such autocorrelation plot.

```
acf(lm1$res)
```



The horizontal dash lines marks the levels autocorrelations to be significantly different from 0.

Remedies to Correct for Autocorrelated Errors: Transformation

Assume the errors ϵ_t 's of the linear model $y_t = \beta_0 + \beta_1 x_t + \epsilon_t$ have the first order autocorrelation AR(1) structure:

$$\epsilon_t = \rho \epsilon_{t-1} + \omega_t, \omega_t \sim N(0, \theta^2)$$

Then

$$y_t - \beta_0 - \beta_1 x_t = \rho(y_{t-1} - \beta_0 - \beta_1 x_{t-1}) + \omega_t$$

$$\underbrace{y_t - \rho y_{t-1}}_{y_t^*} = \beta_0(1 - \rho) + \beta_1 \underbrace{(x_t - \rho x_{t-1})}_{x_t^*} + \omega_t$$

Hence the transformed variables x_t^* and y_t^* satisfy the SLR model:

$$y_t^* = \beta_0^* + \beta_1^* x_t^* + \omega_t$$

, where ω_t are indep. $\sim N(0, \theta^2)$

Interpreting the coefficients

- Transform back!
- The coefficients of the original and the transformed models are related as follows:

$$\beta_0(1 - \rho) = \beta_0^*, \quad \beta_1 = \beta_1^*$$

- However we need to estimate ρ

Cochrane-Orcutt Method

- 1 Fit the OLS model and obtain the residuals e_1, \dots, e_n
- 2 Use the residuals e_1, \dots, e_n to estimate ρ with

$$\hat{\rho} = \frac{\sum_{t=2}^n e_t e_{t-1}}{\sum_{t=1}^n e_t^2}$$

- 3 Compute OLS estimates of β_0^* and β_1^* by regressing $y^* = y_t - \hat{\rho}y_{t-1}$ on $x^* = x_t - \hat{\rho}x_{t-1}$ and use them to find coefficients for the original variables. $\hat{\beta}_0 = \frac{\hat{\beta}_0^*}{1 - \hat{\rho}}$ and $\hat{\beta}_1 = \hat{\beta}_1^*$

- 4 Use the new $\hat{\beta}_0$ and $\hat{\beta}_1$ to calculate the new residuals e_1, \dots, e_n and then go back to Step 2.

- 5 Iterate until the estimates $\hat{\beta}_0$ and $\hat{\beta}_1$ converge.

R code: first iteration

```
x = stock$Stock
y = stock$Expenditure
n = length(y)
fit1 = lm(y ~ x)
res = fit1$res
rho.hat = sum(res[1:(n-1)]*res[2:n]) / sum(res^2 )
rho.hat
```

```
## [1] 0.7506122
```

```
ystar = y[2:n] - rho.hat*y[1:(n-1)]
xstar = x[2:n] - rho.hat*x[1:(n-1)]
fit2 = lm(ystar ~ xstar)
b0.hat = fit2$coef[1]/(1-rho.hat)
b1.hat = fit2$coef[2]
c(b0.hat, b1.hat)
```

```
## (Intercept)      xstar
## -215.310969    2.643443
```


R code: second iteration

```
res = y - b0.hat - b1.hat*x
rho.hat = sum(res[1:(n-1)]*res[2:n]) / sum(res^2 )
rho.hat
```

```
## [1] 0.789962
```

```
ystar = y[2:n] - rho.hat*y[1:(n-1)]
xstar = x[2:n] - rho.hat*x[1:(n-1)]
fit2 = lm(ystar ~ xstar)
b0.hat = fit2$coef[1]/(1-rho.hat)
b1.hat = fit2$coef[2]
c(b0.hat, b1.hat)
```

```
## (Intercept)          xstar
## -225.600207      2.699873
```

R code: More iterations

```
x = stock$Stock
y = stock$Expenditure
n = length(y)
n.iter = 15
rho.iter = vector("numeric", n.iter)
b0.iter = vector("numeric", n.iter)
b1.iter = vector("numeric", n.iter)
fit1 = lm(y ~ x)
res = fit1$res
rho.iter[1] = sum(res[1:(n-1)]*res[2:n]) / sum(res^2 )


for(i in 2:n.iter){
  rho.iter[i] = sum(res[1:(n-1)]*res[2:n]) / sum(res^2 )
  ystar = y[2:n] - rho.iter[i]*y[1:(n-1)]
  xstar = x[2:n] - rho.iter[i]*x[1:(n-1)]
  fit2 = lm(ystar ~ xstar)$coef
  b0.iter[i] = fit2[1]/(1-rho.iter[i])
  b1.iter[i] = fit2[2]
  res = y - b0.iter[i] - b1.iter[i]*x
}
```

R code: More iterations

```
##      rho.iter    b0.iter    b1.iter
## 1  0.7506122      0.0000  0.000000
## 2  0.7506122 -215.3110  2.643443
## 3  0.7899620 -225.6002  2.699873
## 4  0.7977275 -227.8166  2.711905
## 5  0.7995646 -228.3473  2.714777
## 6  0.8000161 -228.4780  2.715485
## 7  0.8001281 -228.5105  2.715660
## 8  0.8001560 -228.5185  2.715704
## 9  0.8001629 -228.5206  2.715715
## 10 0.8001646 -228.5211  2.715717
## 11 0.8001650 -228.5212  2.715718
## 12 0.8001651 -228.5212  2.715718
## 13 0.8001652 -228.5212  2.715718
## 14 0.8001652 -228.5212  2.715718
## 15 0.8001652 -228.5212  2.715718
```

Checking the Independence Assumption After Transformation

```
rho.hat = rho.iter[n.iter]
ystar = y[2:n] - rho.hat*y[1:(n-1)]
xstar = x[2:n] - rho.hat*x[1:(n-1)]
xystar = data.frame(xstar, ystar)
fit2 = lm(ystar ~ xstar)
ggplot(xystar, aes(x=1:(n-1), y = fit2$res)) +
  geom_point() + geom_line() +
  labs(x="Index", y="Residual") + geom_hline(yintercept=0)
```



Runs Test

```
library(tseries)
runs.test(factor(fit2$res > 0)) # two-sided by default
```

```
##
##  Runs Test
##
## data:  factor(fit2$res > 0)
## Standard Normal = -0.69782, p-value = 0.4853
## alternative hypothesis: two.sided
runs.test(factor(fit2$res > 0), alternative = "less")
```

```
##
##  Runs Test
##
## data:  factor(fit2$res > 0)
## Standard Normal = -0.69782, p-value = 0.2426
## alternative hypothesis: less
runs.test(factor(fit2$res > 0), alternative = "greater")
```

```
##
##  Runs Test
##
## data:  factor(fit2$res > 0)
## Standard Normal = -0.69782, p-value = 0.7574
## alternative hypothesis: greater
```

Durbin-Watson Test

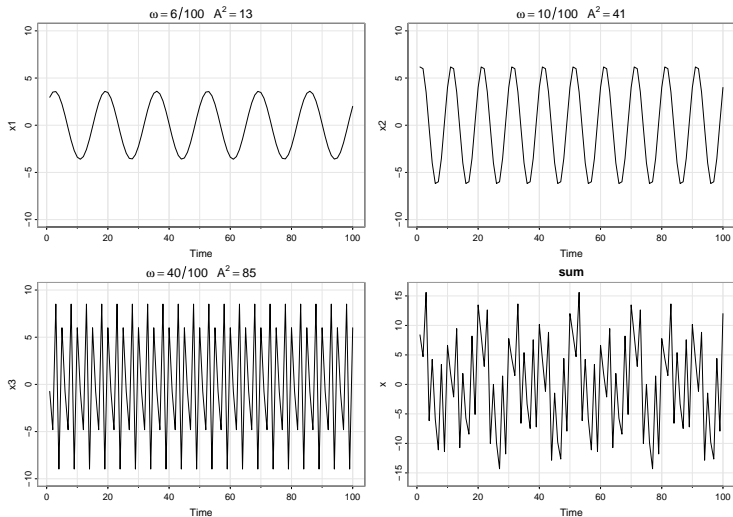
```
library(car)  
durbinWatsonTest(fit2, alt="positive")
```

```
## lag Autocorrelation D-W Statistic p-value  
## 1 0.1824862 1.548671 0.109  
## Alternative hypothesis: rho > 0
```

Autocorrelation and Seasonality

```
x1 = 2*cos(2*pi*1:100*6/100) + 3*sin(2*pi*1:100*6/100)
x2 = 4*cos(2*pi*1:100*10/100) + 5*sin(2*pi*1:100*10/100)
x3 = 6*cos(2*pi*1:100*40/100) + 7*sin(2*pi*1:100*40/100)
x = x1 + x2 + x3
```


Autocorrelation and Seasonality



Autocorrelation

ACF plots can help detect higher-order dependence.

```
acf(x)
```



Durbin-Watson Test

```
fit_syn = lm(x~1)
durbinWatsonTest(fit_syn, alt="positive")
```

```
## lag Autocorrelation D-W Statistic p-value
## 1 -0.1836366 2.336403 0.953
## Alternative hypothesis: rho > 0
```


Durbin-Watson test give large P-values even though there exist significant lag-k autocorrelation (It failed!)

Seasonality (Time and frequency)

- The frequency is measured in cycles per unit time (ω)
- $\omega = 1.0$ the series makes one cycle per time unit
- $\omega = 0.5$ the series makes one cycle every two time units
- In general, for data that occur at discrete time points will need at least two points to determine a cycle, so the highest frequency of interest is 0.5 cycles per point.
- This frequency is called the folding frequency/ Nyquist frequency ($F = f_s/2$)
- Higher frequencies sampled this way will appear at lower frequencies, called aliases


Alias and Folding Frequency ($f=0.4\text{Hz}$)

```
x4 = 2*cos(2*pi*1:100*40/100)
P = abs(fft(x4))^2
Fr = 0:99/100
par(mfrow=c(2,1))
tsplot(x4, ylim=c(-2,2), main = expression(omega==40/100~~~A^2))
tsplot(Fr, P, type="o", xlab="frequency", ylab="Amplitude")
abline(v=.5, lty=2)
```



Alias and Folding Frequency ($f=0.6\text{Hz}$)

```
x4 = 2*cos(2*pi*1:100*60/100)
P = abs(fft(x4))^2
Fr = 0:99/100
par(mfrow=c(2,1))
tsplot(x4, ylim=c(-2,2), main = expression(omega==60/100~~~A^2))
tsplot(Fr, P, type="o", xlab="frequency", ylab="Amplitude")
abline(v=.5, lty=2)
```



Seasonality: The Scaled Periodogram

It may be regarded as a measure of the squared correlation of the data

```
P = abs(2*fft(x)/100)^2  
Fr = 0:99/100  
tsplot(Fr, P, type="o", xlab="frequency", ylab="periodogram")  
abline(v=.5, lty=2)
```

